

Chapter 9: Planning and Learning

Objectives of this chapter:


- ❑ Use of environment models
- ❑ Integration of planning and learning methods

Models

- ❑ **Model**: anything the agent can use to predict how the environment will respond to its actions
- ❑ **Distribution model**: description of all possibilities and their probabilities
 - e.g., $P_{ss'}^a$ and $R_{ss'}^a$, for all s, s' , and $a \in A(s)$
- ❑ **Sample model**: produces sample experiences
 - e.g., a simulation model
- ❑ Both types of models can be used to produce **simulated experience**
- ❑ Often sample models are much easier to come by

Planning

- ❑ **Planning**: any computational process that uses a model to create or improve a policy




- ❑ Planning in AI:

- state-space planning
- plan-space planning (e.g., partial-order planner)

- ❑ We take the following (unusual) view:

- all state-space planning methods involve computing value functions, either explicitly or implicitly
- they all apply backups to simulated experience



Planning Cont.

- ❑ Classical DP methods are state-space planning methods
- ❑ Heuristic search methods are state-space planning methods
- ❑ A planning method based on Q-learning:


Do forever:

1. Select a state, $s \in \mathcal{S}$, and an action, $a \in \mathcal{A}(s)$, at random
2. Send s, a to a sample model, and obtain a sample next state, s' ,
and a sample next reward, r
3. Apply one-step tabular Q-learning to s, a, s', r :
$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

Random-Sample One-Step Tabular Q-Planning

Learning, Planning, and Acting

- Two uses of real experience:
 - **model learning**: to improve the model
 - **direct RL**: to directly improve the value function and policy
- Improving value function and/or policy via a model is sometimes called **indirect RL** or **model-based RL**. Here, we call it **planning**.



Direct vs. Indirect RL

❑ Indirect (model-based) methods:


- make fuller use of experience: get better policy with fewer environment interactions

❑ Direct methods

- simpler
- not affected by bad models

But they are very closely related and can be usefully combined: planning, acting, model learning, and direct RL can occur simultaneously and in parallel

The Dyna Architecture (Sutton 1990)



The Dyna-Q Algorithm

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Do forever:

(a) $s \leftarrow$ current (nonterminal) state

(b) $a \leftarrow \epsilon$ -greedy(s, Q)

(c) Execute action a ; observe resultant state, s' , and reward, r

(d) $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

(e) $Model(s, a) \leftarrow s', r$ (assuming deterministic environment)

(f) Repeat N times:

$s \leftarrow$ random previously observed state

$a \leftarrow$ random action previously taken in s

$s', r \leftarrow Model(s, a)$

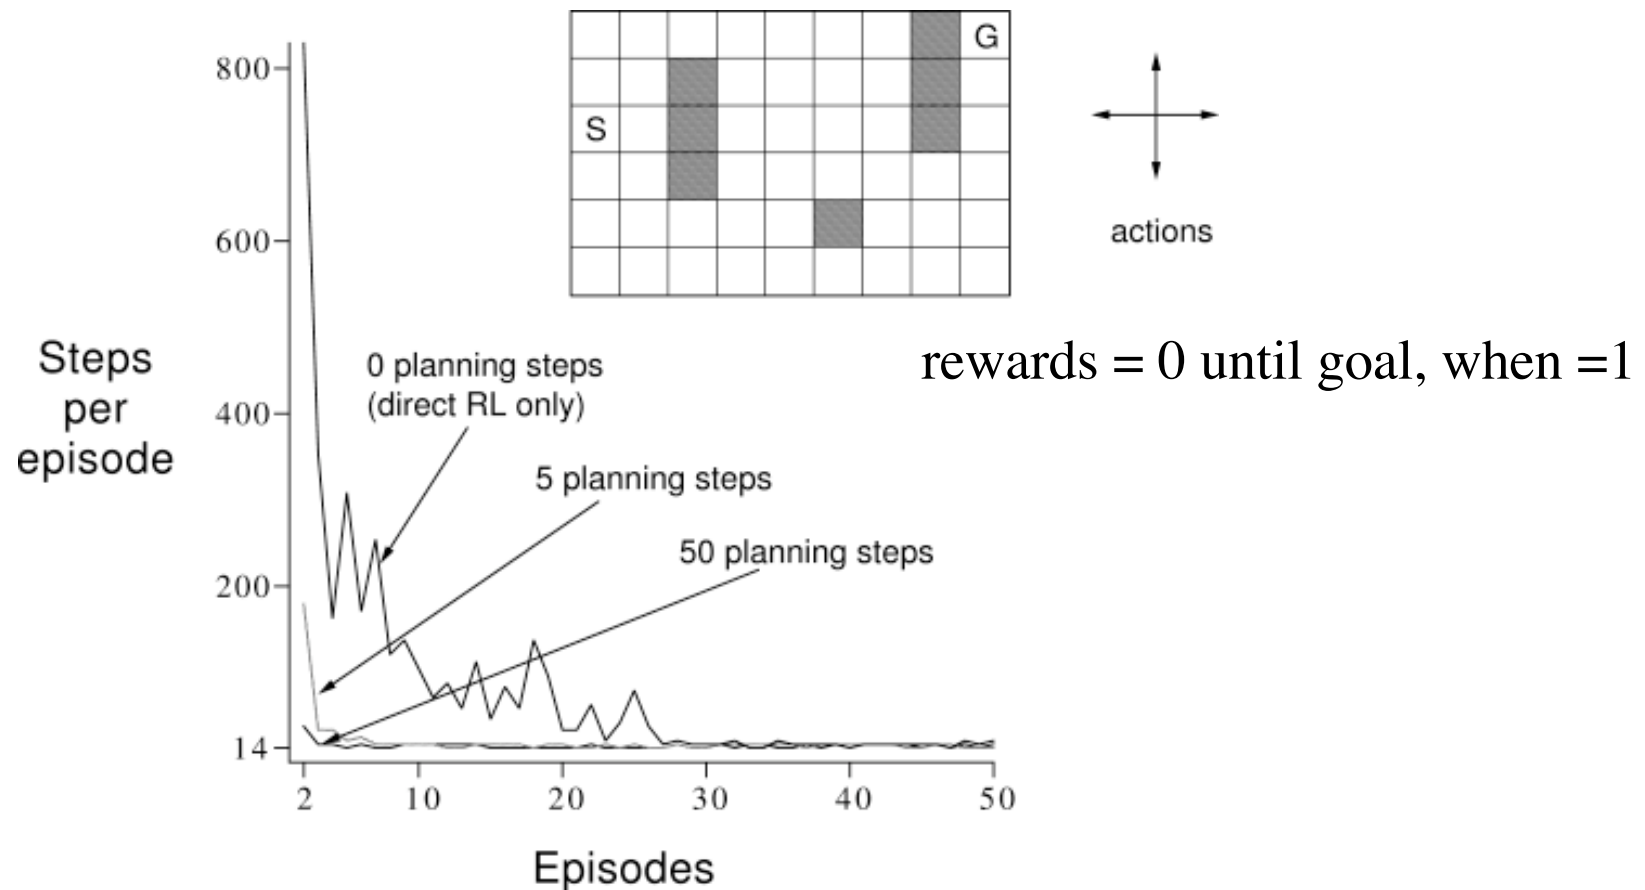
$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

direct
RL

model learning


planning

Dyna-Q on a Simple Maze




Dyna-Q Snapshots: Midway in 2nd Episode

WITHOUT PLANNING ($N=0$)




WITH PLANNING ($N=50$)




When the Model is Wrong: Blocking Maze

The changed environment is harder



Shortcut Maze

The changed environment is easier



What is Dyna-Q⁺?

- ❑ Uses an “exploration bonus”:
 - Keeps track of time since each state-action pair was tried for real
 - An extra reward is added for transitions caused by state-action pairs related to how long ago they were tried: the longer unvisited, the more reward for visiting
 - The agent actually “plans” how to visit long unvisited states

Prioritized Sweeping

- ❑ Which states or state-action pairs should be generated during planning?
- ❑ Work backwards from states whose values have just changed:
 - Maintain a queue of state-action pairs whose values would change a lot if backed up, prioritized by the size of the change
 - When a new backup occurs, insert predecessors according to their priorities
 - Always perform backups from first in queue
- ❑ Moore and Atkeson 1993; Peng and Williams, 1993


Prioritized Sweeping

Initialize $Q(s, a)$, $Model(s, a)$, for all s, a , and $PQueue$ to empty
Do forever:


- (a) $s \leftarrow$ current (nonterminal) state
- (b) $a \leftarrow policy(s, Q)$
- (c) Execute action a ; observe resultant state, s' , and reward, r
- (d) $Model(s, a) \leftarrow s', r$
- (e) $p \leftarrow |r + \gamma \max_{a'} Q(s', a') - Q(s, a)|$.
- (f) if $p > \theta$, then insert s, a into $PQueue$ with priority p
- (g) Repeat N times, while $PQueue$ is not empty:
 - $s, a \leftarrow first(PQueue)$
 - $s', r \leftarrow Model(s, a)$
 - $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
 - Repeat, for all \bar{s}, \bar{a} predicted to lead to s :
 - $\bar{r} \leftarrow$ predicted reward
 - $p \leftarrow |\bar{r} + \gamma \max_a Q(s, a) - Q(\bar{s}, \bar{a})|$.
 - if $p > \theta$ then insert \bar{s}, \bar{a} into $PQueue$ with priority p

Prioritized Sweeping vs. Dyna-Q

Both use $N=5$ backups per
environmental i



Rod Maneuvering (Moore and Atkeson 1993)




Full and Sample (One-Step) Backups

Value
estimated


Full backups
(DP)

Sample backups
(one-step TD)



$V^\pi(s)$





$V^*(s)$




$Q^\pi(a,s)$



$Q^*(a,s)$




Full vs. Sample Backups



b successor states, equally likely; initial error = 1;
assume all next states' values are correct



Trajectory Sampling

- ❑ **Trajectory sampling**: perform backups along simulated trajectories
- ❑ This samples from the on-policy distribution
- ❑ Advantages when function approximation is used
- ❑ Focusing of computation: can cause vast uninteresting parts of the state space to be (usefully) ignored:




Trajectory Sampling Experiment

- ❑ one-step full tabular backups
- ❑ uniform: cycled through all state-action pairs
- ❑ on-policy: backed up along simulated trajectories
- ❑ 200 randomly generated undiscounted episodic tasks
- ❑ 2 actions for each state, each with b equally likely next states
- ❑ .1 prob of transition to terminal state
- ❑ expected reward on each transition selected from mean 0 variance 1 Gaussian



Heuristic Search

- ❑ Used for action selection, not for changing a value function (=heuristic evaluation function)
- ❑ Backed-up values are computed, but typically discarded
- ❑ Extension of the idea of a greedy policy — only deeper
- ❑ Also suggests ways to select states to backup: smart focusing:



Summary

- ❑ Emphasized close relationship between planning and learning
- ❑ Important distinction between **distribution models** and **sample models**
- ❑ Looked at some ways to integrate planning and learning
 - synergy among planning, acting, model learning
- ❑ Distribution of backups: focus of the computation
 - trajectory sampling: backup along trajectories
 - prioritized sweeping
 - heuristic search
- ❑ Size of backups: full vs. sample; deep vs. shallow